

Scalable Inversion of Contests with Correlated Performances, Including Softmax and Multinomial Probit

Peter Cotton*

August 2026

Abstract

Multinomial probit choice probabilities over n alternatives are Gaussian orthant integrals, and for thirty years econometrics has computed them by simulation, most commonly with the GHK sampler, whose R draws cost $O(n^2R)$ per probability, n -fold that for a full vector, with error $O(R^{-1/2})$. We observe that the choice event is one-dimensional: the probability that alternative i wins is an integral over the winning value, and one shared survival field prices all n alternatives at once. For covariances under which performances are conditionally independent given a few shared variables (factor, block, nested and hierarchical structures, the restrictions applied multinomial probit imposes in practice), one shared field computes all n probabilities, and matrix-free Jacobian-vector products, in $O(nLQ)$ operations over a lattice of L winning values and Q factor quadrature nodes, and inverts observed probabilities back to utilities under a fixed, known covariance. At $n = 200$ it is two hundred times faster than GHK at 10^4 draws, with error below the comparison's resolution; GHK's measured cost grows superquadratically (exponent 2.15 to 2.8 across the measured range) against the lattice's near-linear growth, so the ratio widens by orders of magnitude. The same two calls price and invert softmax, Thurstone–Mosteller, factor multinomial probit, and a hierarchy of team and sector models; forward pricing is benchmarked to ten million alternatives.

1 Introduction

A decision maker faces n alternatives and chooses the one with the highest utility. Write $U_i = \mu_i + \varepsilon_i$ with $\varepsilon \sim N(0, \Sigma)$; the probability that alternative i is chosen is

$$p_i = \Pr\{U_i \geq U_j \text{ for all } j\} = \Pr\{\tilde{U} \geq 0\}, \quad \tilde{U}_j = U_i - U_j, \quad (1)$$

an $(n - 1)$ -dimensional Gaussian orthant integral. This is the multinomial probit model. Its appeal has been understood since Thurstone (1927): unlike the logit family it imposes no independence of irrelevant alternatives, because Σ lets alternatives be close substitutes. Its curse has been understood just as long. The integral (1) has no closed form, and estimation needs it, and its derivatives, at every trial value of the parameters.

The workhorse answer is the GHK simulator of Geweke (1989), Börsch-Supan and Hajivassiliou (1993), Keane (1994) and Hajivassiliou et al. (1996). GHK factors Σ by Cholesky and

*peter.cotton@microprediction.com. The reference implementation, its parity-locked ports (Python, R, Rust, JavaScript) and the seeded scripts behind every number in this paper live at github.com/microprediction/winning; this revision's tables were produced at repository tag `paper-r1`.

writes the orthant probability as a nested sequence of one-dimensional truncated normal probabilities, each conditioning on draws for the previous coordinates. Averaging R such importance weights gives an unbiased, smooth estimate of one p_i at cost $O(n^2)$ per draw. The estimator made maximum simulated likelihood practical and it remains the default in textbooks (Train, 2009). Its costs are also standard: error shrinks as $O(R^{-1/2})$; the log of a simulated probability is biased at $O(1/R)$; gradients inherit the simulation noise; and a full probability vector, which the likelihood needs, costs n separate runs.

This paper takes a different route to the same object, in a min-wins convention fixed here once: performances are times, $X_i = -U_i$, and the lowest time wins. The choice event is one-dimensional. Contestant i wins when its time beats the field, so

$$p_i = \int f_i(x) \prod_{j \neq i} S_j(x) dx, \quad (2)$$

where f_i is the density of X_i and S_j the survival of X_j : integrate over the winning time x , requiring every rival to come in slower. For independent performances Cotton (2021) computed (2) for all n on a lattice by building the field $\sum_j \log S_j$ once and dividing each alternative out, and inverted the map from probabilities back to μ . The present paper extends the field construction to correlated performances by conditioning: whenever Σ makes performances independent given a few shared variables, the field factorizes, and all n probabilities, Jacobian–vector products, and each inversion iteration cost $O(nLQ)$ for a lattice of L points and Q quadrature nodes; the Jacobian itself is a weighted graph Laplacian applied as an operator, never materialized. No per-alternative repetition; error controlled by quadrature, deterministic for the Gauss–Hermite rules and reproducibly randomized for scrambled-Sobol nodes.

Only utility contrasts and a normalized scale are identified in multinomial probit, so the raw covariance is overparameterized, and the contrast covariance may in principle carry unrestricted parameters. Applied work nevertheless imposes structured covariance as a matter of course, for parsimony and for computation (Train, 2009); the grammars here cover the structured families in common use. The table lists them: the first six rows are nested exactly and each is priced and inverted by the same two calls; the last row is a fitted approximation developed in §6.

model	covariance grammar	source
Luce / softmax	independent, minimum-Gumbel base, $D = \pi^2/6$	exact identity
Thurstone–Mosteller	independent, normal base	Cotton (2021)
factor multinomial probit	$VV' + \text{diag}(D)$	§4.1
teams, conferences	blocks: rank- r effect per cluster	§4.2
market plus sectors	nested: blocks $\times$ global factor	§4.2
hierarchies	tree: uniform shared effects	§4.3
dense Σ	fitted grammar + residual factors (approx.)	§6

Scale motivates second. The next table reports measured wall clock for all n probabilities (Rust kernels, one laptop) against extrapolations of GHK’s measured cost. The measurements behind the law: the complete-vector GHK timings at $R = 10^3$ are 0.57, 5.66 and 51.3 seconds at $n = 200, 500$ and 1000 (times at $R = 10^4$ scale by ten), a log-log slope of 2.8 over that range, while the small- n table below gives 2.15 from $n = 10$ to 200: the law is superquadratic and steepening. Extrapolations are shown as brackets over the two measured exponents, and every

timing is regenerated by the committed `bench.py`. At these budgets GHK’s total variation error is near 10^{-2} ; the lattice error is near 10^{-8} .

n	lattice, measured	GHK at 10^4 draws, cost-law bracket	multiplier
10^4	0.18 s	20 hours – 3.7 days	$4 \times 10^5 - 2 \times 10^6$
10^5	2.7 s	0.3 – 6.5 years	$4 \times 10^6 - 8 \times 10^7$
10^6	29 s	46 – 4,100 years	$5 \times 10^7 - 4 \times 10^9$

Section 2 places the method among the alternatives. Section 3 gives the independent lattice race; Section 4 extends it across the covariance grammar, factor to blocks to trees, each with its recursion written out. Section 5 gives the Jacobian operator and the inversion. Section 6 reduces dense covariance to the grammar family, approximately, and reports accuracy per named random-correlation ensemble. Section 7 contains the benchmarks.

2 The alternatives

Deterministic quadrature. Genz (1992) transformed the orthant integral to the unit cube and integrated by adaptive and quasi-Monte Carlo rules; the `mvtnorm` implementation is the standard in R. Accuracy is not the issue: in the benchmark below it agrees with the lattice to the reference’s own noise. The shape is the issue. It treats one probability at a time, so the all- n vector at $n = 30$ costs three hundred times the lattice, and the gap widens with n .

Frequency simulation. The crude frequency estimator of Lerman and Manski (1981) counts argmax draws. It is unbiased for the whole vector at once but not smooth in parameters, and its per-entry relative error explodes for small p_i . Measured below: at wall clock matched to the lattice it carries twelve to twenty times the total-variation error, records zero counts on tail alternatives at $n = 30$, and at ten times the budget remains a factor of two to four behind.

GHK. Defined above. Its dominance is documentable rather than folklore. Stata’s `cmmprobit` manual states that the likelihood evaluator “implements the Geweke–Hajivassiliou–Keane (GHK) algorithm to approximate the multivariate distribution function” (StataCorp, 2023; Gates, 2006), and the command’s entire `intmethod` menu (Hammersley, Halton, or pseudo-random point sets) selects only the sequence feeding the same simulator: the market-leading implementation offers no non-GHK method at all. R’s `mlogit` and `mvProbit` default to GHK likewise (Croissant, 2020), and the GHK-based `mvprobit` of Cappellari and Jenkins (2003) carried the method to applied Stata work at large. Two properties matter for comparison: it is smooth in parameters, which frequency counting is not, and it prices one alternative per run. Stern (1992) gave a related decomposition with an analytic leading term.

Bayesian data augmentation. Albert and Chib (1993) and McCulloch and Rossi (1994) sidestep the integral entirely by Gibbs sampling the latent utilities; Imai and van Dyk (2005) implement the marginal data-augmentation version in the MNP package. The cost moves into the Markov chain, and the likelihood surface itself is never available, which forecloses everything that needs it. No pricing benchmark is possible for the same reason: the method never produces the probabilities.

Analytic approximations. Clark (1961) matched moments of pairwise maxima; Mendell and Elston (1974) and Solow (1990) approximate by sequential conditioning on first and second moments. These are fast, sometimes startlingly accurate, and carry no error control, and the benchmark below measures all three properties at once: Mendell–Elston prices the $n = 10$ vector in two milliseconds within 2×10^{-3} of the reference, then drifts to 3×10^{-2} at $n = 30$ with tail probabilities off by a factor of nearly four, and nothing in the output announces the change. The Clark-type recursion behaves alike. The family is current practice, not history: Bhat (2011) builds the MACML estimator on precisely these approximations, and the `Rprobit` package implements it (Bauer, 2024), with its asymptotics examined by Batram and Bauer (2019). Expectation propagation is the modern member of the same family. We benchmark the family through its classical members, which is also a benchmark of MACML’s core ingredient.

Minimax tilting. Botev (2017) computes single Gaussian rectangle probabilities with bounded relative error deep into the tails, in dimensions up to roughly a thousand. It is the accuracy reference, and we use it as the referee for our own tail claims; it prices one probability per run.

Every smooth, likelihood-grade method above prices one probability at a time; frequency counting obtains the vector but not derivatives. Estimation wants all n and their derivatives at every parameter value. Sharing the field across alternatives is the algorithmic point of this paper, and it is what makes the n -fold and the derivative costs disappear.

3 The lattice race

In the min-wins convention of the introduction, let $X_i = \mu_i + \sigma_i \epsilon_i$ with ϵ_i i.i.d. from a standardized base density (normal, Gumbel, or any density supplied as code), and write f_i, S_i for the density and survival of X_i .

The algorithm evaluates (2) on a lattice $x_1 < \dots < x_L$:

1. compute $\log S_j(x_\ell)$ and $\log f_j(x_\ell)$ for all j, ℓ ;
2. accumulate the field $L(x_\ell) = \sum_j \log S_j(x_\ell)$;
3. for each i , form the integrand $f_i(x_\ell) e^{L(x_\ell) - \log S_i(x_\ell)}$, which is contestant i ’s density times everyone else’s survival, with i divided out of the shared field in log space;
4. sum times the lattice spacing, and normalize the vector.

The cost is $O(nL)$: the field is built once, not once per contestant. Log-space arithmetic keeps hopeless contestants finite (their $e^{L - \log S_i}$ underflows gracefully rather than dividing zero by zero), and the exponent is clipped at -745 .

The lattice itself is placed on the winner’s bulk, not the ability span. A hopeless contestant wins only by running a winner-class time, so only the region where the winning time lives matters. Bisect the conservative winner c.d.f. envelope $G(x) = 1 - \prod_j S_j(x)$ to $[G^{-1}(\delta), G^{-1}(1 - \delta)]$ and pad two standard deviations, which makes the truncated tail mass negligible against the accuracy floors reported below. Measured: 33 points on this window reach accuracy that a 500-point ability-span lattice cannot, because the span window’s own truncation floors near 6×10^{-11} ; on fields with many hopeless contestants the bulk window is one hundred times more accurate at equal budget.

Two byproducts are cheap once the field is built. Any specified removal counterfactual, the probability j wins after deleting i , costs one more division out of the field, and any specified photo-finish tie density is one more lattice sum; §5 identifies the latter with the off-diagonal Jacobian. The full $n \times n$ matrices of such queries are $O(n^2)$ outputs and are priced per query, not free.

4 The covariance grammar

Three structures make contestants conditionally independent given a few shared variables, in increasing order of hierarchy: a global factor, factors per cluster, and a tree of uniform shared effects. Each keeps the shared field and the $O(nLQ)$ bill; each is exact, not fitted. The approximate reduction of a dense covariance to this family is a separate matter and has its own section (§6).

4.1 Factor

Let $X = \mu + Vf + \text{diag}(\sqrt{D})\epsilon$ with $f \sim N(0, I_k)$ independent of ϵ , so $\Sigma = VV' + \text{diag}(D)$. Conditional on f the performances are independent with means $\mu + Vf$, and

$$p = \sum_q w_q p^{\text{ind}}(\mu + Vf_q), \quad (3)$$

a Q -node quadrature over runs of the independent algorithm, cost $O(nLQ)$. In English: average the independent race over the shared luck. Gauss–Hermite nodes serve $k \leq 2$; scrambled Sobol serves $k \geq 3$.

One default matters. When $D_i \ll \|V_i\|^2$ the conditional race is nearly deterministic and the integrand over f is nearly a step function; a fixed 15-node rule then silently loses up to five percent of total variation, and the loss is invisible to lattice refinement because it is points-invariant. Randomized invariant testing against a QMC referee found this. The default order now scales with the sharpness $\max_i \|V_i\|/\sqrt{D_i}$, capped by rank.

The same pass returns the own-parameter slopes $\partial p_i/\partial \mu_i$ at no extra cost (negative in this min-wins convention: raising a time mean lowers the win probability), which §5 uses as the Newton preconditioner.

4.2 Blocks and nested

Assign contestant i to cluster $c(i)$ with loading $v_i \in \mathbb{R}^r$ and let

$$X_i = \mu_i + v_i' a_{c(i)} + \sqrt{D_i} \epsilon_i, \quad a_c \text{ i.i.d. } N(0, I_r), \quad (4)$$

a block-diagonal rank- r -plus-diagonal covariance: teams, conferences, sibling products. Two independences do the work. Conditional on its own cluster's effect a contestant is independent of its teammates; and distinct clusters are independent outright. So the joint survival factorizes across clusters,

$$\Pr\{X_j > x \text{ for every } j\} = \prod_c G_c(x), \quad G_c(x) = \mathbb{E}_a \prod_{j \in c} S_j(x | a) = \sum_q w_q e^{S_c(x, a_q)}, \quad (5)$$

with $S_c(x, a) = \sum_{j \in c} \log S_j(x | a)$ the cluster's conditional log-survival. In English: $G_c(x)$ is the probability that every member of cluster c survives past x , with the cluster's shared luck integrated out.

Contestant i then wins against a cavity field:

$$p_i = \int h_i(x) \prod_{c \neq c(i)} G_c(x) dx, \quad h_i(x) = \sum_q w_q f_i(x | a_q) e^{S_{c(i)}(x, a_q) - \log S_i(x | a_q)}. \quad (6)$$

The term h_i couples i 's win density to its own teammates within the shared draw of $a_{c(i)}$; the cavity product supplies the rest of the field. The cost is $O(nLQ_a)$ whatever the number of clusters, because the per-cluster factors are accumulated once. Rank two is special: the conditional integrand is smooth and a tensor Gauss–Hermite rule beats scrambled Sobol at equal nodes by a wide measured margin.

Nested. Add one global factor with couplings g_i and a dial $\gamma \in [0, 1]$: conditional on the global draw f_q the model is a block race at shifted abilities, so $p = \sum_q w_q p_{\text{block}}(\mu + \gamma g f_q)$, a finite mixture of block races. At $\gamma = 0$ the clusters are isolated, at $\gamma = 1$ fully coupled, and the covariance contribution scales as $\gamma^2 g g'$.

4.3 Trees

Let a rooted tree have the leaf clusters at its bottom and internal nodes t above, and give node t a scalar effect $b_t \sim N(0, 1)$ applied with uniform strength λ_t to every leaf beneath it:

$$X_i = \mu_i + v_i a_{c(i)} + \sum_{t \in \text{anc}(c(i))} \lambda_t b_t + \sqrt{D_i} \epsilon_i. \quad (7)$$

The implied covariance is hierarchical: two contestants share the variance of exactly their common ancestors. The root's effect is common to every contestant, hence choice-irrelevant, and is fixed to zero. The uniform strength is the crucial restriction, because a shared effect that hits every subtree member equally moves the whole subtree's field by a translation of its argument. Messages therefore remain functions of one variable. Per-leaf loadings on internal effects would force two-dimensional message tables.

The algorithm is two passes over the same lattice.

Upward. Each leaf cluster carries its block factor $G_c(x) = \sum_q w_q e^{S_c(x, a_q)}$ from (5). Each internal node, visited deepest first, combines its children under its own effect:

$$G_t(x) = \sum_q w_q \prod_{c \in \text{children}(t)} G_c(x + \lambda_t a_q). \quad (8)$$

In English: $G_t(x)$ is the probability that every leaf below t survives past x , with t 's effect integrated by quadrature and each child's message shifted because the effect moves that child's whole subtree at once. Shifted evaluations are linear interpolations on the lattice.

Downward. The root’s cavity is one. Each node, visited shallowest first, receives the field of everything outside its subtree:

$$R_t(x) = \left[\sum_q w_q R_{\text{parent}(t)}(x - \lambda_{\text{parent}(t)} a_q) \right] \prod_{s \in \text{siblings}(t)} G_s(x). \quad (9)$$

Contestant i then wins against its own leaf cavity, $p_i = \int h_i(x) R_{c(i)}(x) dx$ with h_i as in (6). The cost is $O(nLQ)$ plus $O(LQ)$ per tree node.

A depth-one tree with zero strengths reproduces the block race to machine precision, which is tested. Against a 2^{22} -path common-random-numbers referee the recursion sits at the simulation noise floor on every resolvable entry at depths one through four (root-mean-square z between 0.75 and 1.1). Entries below the referee’s resolution agree with minimax tilting (Botev, 2017) to 0.3 percent relative at probabilities down to 7×10^{-8} , which is within the tilting estimate’s own error. Given adequate factor quadrature, tail accuracy is carried by the lattice, and the grammar kernels price tails that simulation does not resolve.

Remark 1 (A connection to financial practice). *Long-only portfolio weights could be considered as race probabilities, and there is a tree-race covariance under which hierarchical risk parity’s bisection geometry (López de Prado, 2016) is exactly a race: give internal node t strength $\lambda_t^2 = \rho_t - \rho_{\text{parent}(t)}$ with $\rho_t = \max(1 - 2h_t^2, 0)$ for merge height h_t , and leaf i idiosyncratic variance $1 - \rho_{\text{first}(i)}$; the implied correlation equals the floored cophenetic matrix. The floor is forced, because uniform shared effects cannot represent negative dependence, and omitting it silently inflates other implied correlations above one. The identity turns dendrogram-consistent concentration limits into contest inversions, but portfolio applications are not this paper’s subject.*

5 Jacobians, tie densities, inversion

Proposition 1 (Off-diagonals are photo-finish densities). *For the race (2) and $j \neq i$,*

$$\frac{\partial p_i}{\partial \mu_j} = \int f_i(x) f_j(x) \prod_{k \neq i,j} S_k(x) dx \geq 0, \quad (10)$$

and each row of the Jacobian sums to zero.

The zero row sum says a common shift moves nothing. The off-diagonal is the density of a dead heat between i and j ahead of the field. The Jacobian is a negative weighted graph Laplacian and is applied as an operator: individual entries and Jacobian–vector products are byproducts of the same field pass (per quadrature node a Gram structure over the lattice), while materializing all n^2 entries is a choice with $\Omega(n^2)$ cost that nothing downstream requires. Three derivative objects should be distinguished: the continuum identity (10); its quadrature evaluation; and the exact derivative of the normalized, adaptive-lattice numerical map, which differs from the first two through the normalization quotient and grid motion. The implementation exposes the frozen-grid form, which matches finite differences of the returned map at every lattice resolution, alongside the continuum form, whose symmetry serves preconditioning. The stable pair factorization is $[f_i e^{L - \log S_i}] \times [f_j / S_j]$: the first factor is bounded, and the hazard f_j / S_j grows linearly in the normal base and exponentially in the Gumbel base, either of which the exponent clip absorbs. The symmetric square-root split overflows where survivals vanish.

The block Jacobian is exact, with the within-cluster term computed under the cavity; the nested Jacobian is the mixture of block Jacobians; the tree Jacobian is exact within a cluster and Gram-approximate across clusters, which suffices for Newton because residuals always use the exact forward map. Feasibility-critical callers fall back to finite differences.

The inversion rests on a convex program.

Theorem 1 (The inverse exists and is unique on contrasts). *Fix a covariance of full support (any member of the grammar with every $D_i > 0$) and let $W(\mu) = \mathbb{E} \min_i (\mu_i + \epsilon_i)$, $\epsilon \sim N(0, \Sigma)$. Then W is concave and differentiable with $\nabla W(\mu) = p(\mu)$, and its Hessian is $-L(\mu)$, the negative of the graph Laplacian whose edge weights are the tie densities (10). Every tie density is strictly positive, so W is strictly concave on the contrast space $\mathbf{1}^\perp$. For any target with $p_i^* > 0$ and $\sum_i p_i^* = 1$, the program*

$$\max_{\mu} W(\mu) - \langle p^*, \mu \rangle$$

is constant along $\mathbf{1}$, strictly concave and coercive on $\mathbf{1}^\perp$, and its unique mean-zero maximiser is the inverse: $p(\mu^) = p^*$. If some $p_i^* = 0$ the supremum is approached only as $\mu_i \rightarrow +\infty$; no finite inverse exists, and a floored target recovers a lower bound on that μ_i .*

Proof. A minimum of affine functions of μ is concave and expectation preserves this. Ties have measure zero under full support, so the minimiser is almost surely unique and Danskin’s theorem gives $\partial W / \partial \mu_i = \Pr\{i \text{ wins}\}$. Differentiating once more yields the Jacobian above: off-diagonals (10) and zero row sums, that is $-L(\mu)$. Full support makes every pairwise tie density positive, the tie graph complete, and $L \succ 0$ on $\mathbf{1}^\perp$. Translation invariance $W(\mu + c\mathbf{1}) = W(\mu) + c$ makes the objective constant along $\mathbf{1}$ exactly when the target sums to one. For coercivity take a unit contrast d : the derivative of $W(td) - t\langle p^*, d \rangle$ in t is $\langle p(td) - p^*, d \rangle$, and as $t \rightarrow \infty$ the race concentrates on $\arg\min_i d_i$, so the derivative tends to $\min_i d_i - \langle p^*, d \rangle$, which is strictly negative for interior p^* and nonconstant d . A strictly concave coercive function attains a unique maximum, where the gradient vanishes. $\square$

The program is the Legendre dual of the social surplus function of McFadden (1981), and share inversion as a convex problem is established territory: Berry (1994) inverts market shares in logit-family demand, Galichon and Salanié (2022) identify the general random-utility inversion with optimal transport, and Chiong et al. (2016) compute it by linear programming in the discrete case. What the field representation adds is not the convex program but its derivatives: the Hessian is the tie-density Laplacian of the same field pass, so gauge-fixed Newton runs matrix-free at $O(nLQ)$ per iteration.

Identification lives on the same contrast space. Choice probabilities are invariant to a common utility shock, so the model depends on Σ only through $P\Sigma P$ with $P = I - \mathbf{1}\mathbf{1}^\top/n$: a factor proportional to $\mathbf{1}$, or a tree-root effect applied to every leaf, is choice-irrelevant. The fitting target for a dense covariance is therefore the projected matrix, and the reference implementation provides that fit directly (`factor_model_contrast`); a raw-space fit can spend scarce rank reproducing an irrelevant common component. The introduction’s identified-class sentence is to be read in this sense: structured covariance is the estimable and computable restriction on $\mathbf{1}^\perp$, not a claim that structure exhausts identification.

In practice: an adaptive damped fixed point on the log residual carries the iterate into Newton’s basin, then Newton steps on $\log p(\mu) - \log p^*$, preconditioned by the own-parameter

slopes and solved matrix-free, reach 10^{-8} , which Theorem 1 makes the solution of a well-posed problem; targets below the forward map’s resolution are floored, with the one-sided bound direction the theorem states. Ignoring known structure at inversion time is not a small error: inverting block-generated probabilities under an assumed independent model mislocates abilities by up to three tenths of the field’s spread (the pinned configuration: $n = 200$, thirty clusters, within-cluster correlation 0.65; maximum mislocation 2.49 against a spread of 8.24, median 0.34).

6 Dense covariance

The fitting procedure, stated so it can be reproduced. Standardize Σ to a correlation matrix C . (i) Global factors: the top- k eigenpairs of C , $k = 3$ throughout. (ii) Blocks: average-linkage clustering on $\sqrt{(1 - C)/2}$ cut to 20 clusters; for each cluster, the leading eigenvector of the within-cluster residual $C - VV'$ with its diagonal zeroed, giving one rank-one loading per cluster (clusters are disjoint). Stages (i) and (ii) alternate for eight sweeps, each stage refit with the diagonal excluded so that factors never chase variance owed to the idiosyncratic term. (iii) Residual promotion: the top $m = 5$ eigencomponents of what remains, diagonal zeroed, appended as further factor columns; negative eigenvalues are not promoted. (iv) The idiosyncratic variances are $D_i = \max(C_{ii} - \sum_r V_{ir}^2, 10^{-3})$, so the assembled model is positive semidefinite by construction as a sum of Gram terms and a nonnegative diagonal. The loss at every stage is Frobenius error on the off-diagonal entries. The constants k , the cluster count, the sweep count and m are fixed, not tuned per ensemble; a spectral rule for m was tested and rejected (it saturates its cap for marginal gains and does not repair the failure case below). What results is a factor-plus-blocks race priced in one deterministic call. At $n = 2000$ the call takes five seconds against thirty-six for a 10^6 -draw simulation, agrees with it at its noise floor, and prices the 82 percent of tail alternatives for which the simulation records zero wins.

The fit is performed in raw correlation space. Because choices depend on Σ only through the contrast projection $P\Sigma P$, we also ran the identical pipeline with stage (i) replaced by the projected fit (`factor_model_contrast`): on these ensembles it is indistinguishable on most rows and two to six times worse on several (factor with sparse links, hierarchical, block equicorrelation, kernel). The mechanism is instructive. Projection pays off when Σ carries a large choice-irrelevant common component, and none of these correlation ensembles does; meanwhile several truths lie in or near the raw rank-plus-diagonal class, and $P\Sigma P$ of such a matrix no longer does, so projecting first manufactures model error where the raw fit had none. The projected fit remains the right tool when a dominant common shock is present.

Accuracy depends on the generating ensemble, so it is reported per named ensemble of the `randomcov` library (one draw per ensemble at $n = 300$, seed fixed, judged on entries a two-million-path referee resolves; the complete stratified table, all fifteen ensembles by probability band with $\frac{1}{2}\|\cdot\|_1$ and $\|\cdot\|_\infty$, is generated by the committed `run_ensembles3.py`). Representative rows, absolute error:

ensemble	$\frac{1}{2}\ p - \hat{p}\ _1$	$\ p - \hat{p}\ _\infty$	med, $p > 10^{-2}$	med, 10^{-3} – 10^{-2}	med, $< 10^{-3}$
factor + sparse links	2.8×10^{-3}	7.1×10^{-4}	1.2×10^{-4}	2.8×10^{-5}	6.5×10^{-6}
sparse-precision graphical	4.8×10^{-3}	5.7×10^{-4}	3.2×10^{-4}	5.1×10^{-5}	5.8×10^{-6}
Wishart	1.2×10^{-2}	3.2×10^{-3}	5.5×10^{-4}	1.1×10^{-4}	1.1×10^{-5}
residuals (dense-strong)	6.0×10^{-2}	1.7×10^{-2}	2.3×10^{-3}	6.8×10^{-4}	5.8×10^{-5}
kernel (Matérn/RBF)	3.8×10^{-2}	1.3×10^{-2}	8.8×10^{-3}	4.3×10^{-3}	4.4×10^{-4}

Eleven further ensembles (Marchenko–Pastur and spiked spectra, LKJ, onion, vine, hierarchical, block equicorrelation, AR(1), elliptope walks, and two others) fall between the first and fourth rows. The kernel row is the failure case. Kernel covariance is locality, which a low-rank factorization serves badly. The Matérn family additionally admits a sparse-precision representation through the SPDE connection (Lindgren et al., 2011); the squared-exponential does not share that route, and its spectrum decays faster than any polynomial, so the two cases likely want different specialized methods. Neither offers a low-rank fit nothing to hold. The right algorithm there is sequential conditioning with cheap sparse-precision draws, and a production front door should dispatch on the structure it detects.

7 Benchmarks

Against the field. All- n choice probabilities under the rank-2 factor covariance family of Section 2, against a 2^{20} -point QMC reference. Tail is the maximum absolute log-error over reference probabilities in $[10^{-4}, 10^{-3}]$; the $n = 10$ draw has none. Frequency simulation is run at wall clock matched to the lattice and at ten times that budget; “zeros” counts tail alternatives that received no draws.

method	$n = 10$		$n = 30$		tail
	time	TV	time	TV	
lattice race	2.5 ms	1.8×10^{-4}	4.7 ms	8.3×10^{-4}	0.07
Genz, full vector	46 ms	1.8×10^{-4}	1.40 s	8.3×10^{-4}	0.07
frequency, matched time	—	3.8×10^{-3}	—	1.0×10^{-2}	0.45, 2 zeros
frequency, $10\times$ time	—	9.2×10^{-4}	—	2.0×10^{-3}	0.21
Mendell–Elston	2.0 ms	1.8×10^{-3}	18 ms	3.2×10^{-2}	1.30
Clark-type	1.4 ms	1.1×10^{-2}	17 ms	1.6×10^{-2}	1.33

Genz and the lattice agree to the reference’s noise and to each other; the $300\times$ cost ratio at $n = 30$ is the one-at-a-time shape, not the integrator. At the paper’s headline scale the shape is decisive: one probability at $n = 200$ costs 4.6 s (values agreeing with the lattice to six digits), so the full vector costs fifteen minutes against 26 ms for the lattice, a factor of 3×10^4 . The sequential-conditioning family is the only entry whose error grows silently: its tail entries are off by $e^{1.3} \approx 3.7$ with nothing in the output to say so. The race’s tail figure is the reference’s own noise; lattice self-convergence places its error near 10^{-8} .

Referee agreement. A replay harness maps each benchmark probability to its difference or-thant and prices it through the two field-standard R implementations. `mvtnorm`’s Genz–Bretz agrees with the lattice to within its own reported error bound in every family: 3.8×10^{-7} against a bound of 8.1×10^{-7} at $n = 10$, and 2.7×10^{-8} against 5.8×10^{-8} on a case whose smallest probability is 3×10^{-10} . Botev’s minimax tilting agrees to its own Monte Carlo noise, which

in relative terms is 2.2×10^{-4} at $p \approx 3 \times 10^{-10}$: the tail claim rests on the named accuracy reference, not only on self-convergence. On the independent family an adaptive-quadrature referee reaches machine precision and the lattice matches it to 1.9×10^{-15} . An invariance battery (two-runner closed form, translation, permutation, symmetry) holds to 10^{-16} throughout, except the two-runner comparison at 7×10^{-9} , which is the lattice discretization itself.

Stress boundaries. An adversarial battery pushes each failure axis until something gives. Depth: against pure-relative adaptive quadrature, relative error is 4×10^{-13} at $p \approx 10^{-30}$ and 2×10^{-8} at $p \approx 10^{-50}$, with the first genuine breakdown near $p \approx 10^{-119}$, a laggard twenty-five standard deviations behind. Inversion: the round trip $p \rightarrow \mu \rightarrow p$ holds 10^{-9} in log even when the target contains a 10^{-10} entry or ties at the 10^{-9} level. The Gumbel base reproduces softmax to 10^{-15} at $n = 1000$; exact duplicates split exactly, and an ϵ perturbation moves the split linearly. The one operative boundary is factor strength. When loadings make implied correlations exceed roughly 0.9, the argmax integrand loses smoothness in factor space and Gauss–Hermite converges slowly in the node count: at loading scale 4 the 25-node rule still sits at 8×10^{-3} total variation. Scrambled-Sobol factor nodes restore reference-level accuracy (4×10^{-4}) at the same call: heavy correlation wants quasi-Monte-Carlo nodes, not deeper polynomial rules. The default now escalates the family automatically past sharpness 3, in both the Python and R implementations (scrambled Sobol and Halton respectively), which agree to 5×10^{-4} on the worst case; the figures above are the pre-escalation diagnosis.

Against GHK. All- n choice probabilities under a rank-2 factor covariance, against a 2^{20} -point QMC reference. TV is a bulk metric and blind to the tails, so the same tail column as above is reported: the maximum absolute log-error over reference probabilities in $[10^{-4}, 10^{-3}]$, which for probabilities this small is the log-odds error, and which the reference’s own counting noise floors near 0.1.

n	lattice race			GHK, $R = 10^3$			GHK, $R = 10^4$		
	time	TV	tail	time	TV	tail	time	TV	tail
10	2.3 ms	1.8×10^{-4}	—	1.0 ms	8.4×10^{-3}	—	8 ms	2.1×10^{-3}	—
50	7.0 ms	1.7×10^{-3}	0.11	21 ms	4.2×10^{-2}	1.10	214 ms	1.2×10^{-2}	0.29
200	25 ms	2.4×10^{-3}	0.15	436 ms	3.2×10^{-2}	1.50	4.8 s	1.2×10^{-2}	0.35

The $n = 10$ draw has no reference probability in the tail band. The lattice’s tail figures equal the reference’s own noise; GHK’s are two to ten times larger in log terms even at ten times the draws, which is the tail blindness the TV column cannot show.

The lattice bounds are the reference’s own noise; self-convergence places the lattice error near 10^{-8} . GHK error contracts as $O(R^{-1/2})$ and its measured complete-vector cost grows superquadratically, with exponent 2.15 to 2.8 across the measured range. The lattice is deterministic, supplies smooth matrix-free derivatives free of simulation noise (simulated likelihoods can be differentiated, but inherit the noise), and prices the tails. Applied multinomial probit imposes structured covariance for parsimony and for estimability of the contrast covariance, and for those structured families we consider GHK replaced. GHK retains general rectangle probabilities and locality-structured covariance, and Botev (2017) remains the reference for single extreme tail probabilities.

Estimation. The pricing advantage transfers to estimation. Simulate $T = 50,000$ choices among $n = 30$ alternatives under a known rank-2 factor covariance and estimate the mean-zero utilities by maximum likelihood, exact probabilities and analytic score against GHK probabilities with common random numbers and finite-difference gradients, eight replications:

estimator	RMSE over identified alternatives	median fit time
exact gradient MLE	0.0187	4.9 s
MSL, $R = 100$	0.0305	4.4 s
MSL, $R = 1000$	0.0198	40.7 s

At equal fitting time the simulated estimator pays a 63 percent error premium; at eight times the budget it remains behind. The exact estimator sits at the statistical floor. Deterministic quadrature does not rescue estimation either: at $n = 30$ one probability vector by Genz–Bretz costs 1.4 s and supplies no analytic score, so a single finite-difference gradient step (31 vector evaluations) costs roughly eight times the entire exact-gradient fit.

Scale. Ten million contestants price in 245 seconds under the block grammar (ten thousand clusters, 257-point lattice, the streaming field pass; the committed `bench.py tenmillion`), and the Rust classic-calibration kernel runs 232 times the speed of the pure-Python implementation. Throughput at this size is a systems statement; the accuracy evidence lives at the scales the referees above can reach.

Parity. One vector file embeds the inputs and outputs of 22 scenarios spanning every claim above. The four language implementations of §8 replay them: twenty at machine precision, the optimizer-mediated scenarios to 10^{-6} .

8 Availability

`pip install winning` is pure Python; `winning[fast]` adds the Rust kernels as one abi3 wheel per platform. A base-R package mirrors the full interface, and a zero-dependency browser port drives a live demonstration at winning.microprediction.org.

References

- J. H. Albert, S. Chib. Bayesian analysis of binary and polychotomous response data. *JASA* 88:669–679, 1993.
- S. Berry. Estimating discrete-choice models of product differentiation. *RAND Journal of Economics* 25(2):242–262, 1994.
- Z. I. Botev. The normal law under linear restrictions: simulation and estimation via minimax tilting. *Journal of the Royal Statistical Society B* 79(1):125–148, 2017.
- A. Börsch-Supan, V. Hajivassiliou. Smooth unbiased multivariate probability simulators for maximum likelihood estimation of limited dependent variable models. *Journal of Econometrics* 58:347–368, 1993.

- M. Batram, D. Bauer. On consistency of the MACML approach to discrete choice modelling. *Journal of Choice Modelling* 30:1–16, 2019.
- D. Bauer. *Rprobit: R package for the estimation of multinomial probit models*. github.com/dbauer72/Rprobit.
- C. R. Bhat. The maximum approximate composite marginal likelihood (MACML) estimation of multinomial probit-based unordered response choice models. *Transportation Research Part B* 45(7):923–939, 2011.
- L. Cappellari, S. P. Jenkins. Multivariate probit regression using simulated maximum likelihood. *Stata Journal* 3(3):278–294, 2003.
- K. Chiong, A. Galichon, M. Shum. Duality in dynamic discrete-choice models. *Quantitative Economics* 7(1):83–115, 2016.
- C. E. Clark. The greatest of a finite set of random variables. *Operations Research* 9:145–162, 1961.
- Y. Croissant. Estimation of random utility models in R: the mlogit package. *Journal of Statistical Software* 95(11):1–41, 2020.
- P. Cotton. Inferring relative ability from winning probability in multi-entrant contests. *SIAM Journal on Financial Mathematics* 12(1):295–317, 2021.
- A. Galichon, B. Salanié. Cupid’s invisible hand: social surplus and identification in matching models. *Review of Economic Studies* 89(5):2600–2629, 2022.
- R. Gates. A Mata Geweke–Hajivassiliou–Keane multivariate normal simulator. *Stata Journal* 6(2):190–213, 2006.
- A. Genz. Numerical computation of multivariate normal probabilities. *Journal of Computational and Graphical Statistics* 1:141–149, 1992.
- J. Geweke. Bayesian inference in econometric models using Monte Carlo integration. *Econometrica* 57:1317–1339, 1989.
- V. Hajivassiliou, D. McFadden, P. Ruud. Simulation of multivariate normal rectangle probabilities and their derivatives. *Journal of Econometrics* 72:85–134, 1996.
- K. Imai, D. A. van Dyk. MNP: R package for fitting the multinomial probit model. *Journal of Statistical Software* 14(3):1–32, 2005.
- M. P. Keane. A computationally practical simulation estimator for panel data. *Econometrica* 62:95–116, 1994.
- S. Lerman, C. Manski. On the use of simulated frequencies to approximate choice probabilities. In *Structural Analysis of Discrete Data*, MIT Press, 1981.
- F. Lindgren, H. Rue, J. Lindström. An explicit link between Gaussian fields and Gaussian Markov random fields. *Journal of the Royal Statistical Society B* 73(4):423–498, 2011.
- M. López de Prado. Building diversified portfolios that outperform out of sample. *Journal of Portfolio Management* 42(4):59–69, 2016.
- D. McFadden. Econometric models of probabilistic choice. In C. Manski, D. McFadden, eds., *Structural Analysis of Discrete Data*, MIT Press, 1981.

- R. McCulloch, P. Rossi. An exact likelihood analysis of the multinomial probit model. *Journal of Econometrics* 64:207–240, 1994.
- N. Mendell, R. Elston. Multifactorial qualitative traits: genetic analysis and prediction of recurrence risks. *Biometrics* 30:41–57, 1974.
- A. R. Solow. A method for approximating multivariate normal orthant probabilities. *Journal of Statistical Computation and Simulation* 37:225–229, 1990.
- S. Stern. A method for smoothing simulated moments of discrete probabilities in multinomial probit models. *Econometrica* 60:943–952, 1992.
- StataCorp. *Stata Choice Models Reference Manual: cmmprobit, asmprobit*. Stata Press, College Station, 2023.
- L. L. Thurstone. A law of comparative judgment. *Psychological Review* 34:273–286, 1927.
- K. Train. *Discrete Choice Methods with Simulation*. 2nd ed., Cambridge, 2009.